

Optimization of Search Performance in BreastCare Trial using Trulens

Overview

In this report, I aim to measure and optimize the search performance of my project by focusing on key metrics such as:

- **Context Relevance**
- **Latency**
- **Total Cost**

To achieve this, I conducted multiple experiments involving adjustments to prompt engineering, chunk sizes, chunk overlap, and retrieval limits. The experiments were conducted using the following question:

"What is the D3L-001 trial, and what outcomes does it aim to achieve?"

Experiment 1

Settings:

- Chunk size: 2000
- Chunk overlap: 800

```
12
13 class text_chunker:
14
15     def process(self, pdf_text: str):
16
17         text_splitter = RecursiveCharacterTextSplitter(
18             chunk_size = 2000, #Adjust this as you see fit
19             chunk_overlap = 800, #This let's text have some form of overlap. Useful for keeping chunks contextual
20             length_function = len
21         )
```

Prompt Used:

```
16
17 @instrument
18 def generate_completion(self, query: str, context_str: list) -> str:
19     """
20     Generate answer from context.
21     """
22     prompt = f"""
23     You are an intelligent assistant specialized in breast cancer clinical trials.
24     Your responses should focus on trial information, eligibility requirements, and next steps.
25     <user_query>{query}</user_query>
26     Context: {context_str}
27     Question:
28     {query}
29     Answer:
30     """
31     return Complete("mistral-large2", prompt)
```

Results: In this configuration, the metrics (context relevance, latency, and cost) were not optimal. The app's performance needed improvement, prompting further experimentation.

Python v as cell38

```
1 session.get_leaderboard()
```

app_name	app_version	Answer Relevance	Context Relevance	latency	total_cost
BreastCareTrial	simple	1	0.6667	18.5886	4.1817
BreastCareTrial Retriever	base	None	1	0.4598	0

Experiment 2:

Settings:

- Chunk size: 1500
- Chunk overlap: 400

```
12
13 class text_chunker:
14
15     def process(self, pdf_text: str):
16         # Adjusted chunk size and overlap
17         text_splitter = RecursiveCharacterTextSplitter(
18             chunk_size = 1500, # Lower chunk size for better latency
19             chunk_overlap = 400, # Moderate overlap to maintain context
20             length_function = len
21         )
22
23         chunks = text_splitter.split_text(pdf_text)
24         df = pd.DataFrame(chunks, columns=['chunks'])
25
```

Prompt:

```
20         Generate answer from context.
21         """
22         prompt = f"""
23         You are an intelligent assistant specialized in breast cancer clinical trials.
24         Your responses should focus on trial information, eligibility requirements, and next steps.
25         <user_query>{query}</user_query>
26         Context: {context_str}
27         Question:
28         {query}
29         Answer:
30         """
31         return Complete("mistral-large2", prompt)
32
33     @instrument
34     def query(self, query: str) -> str:
35         context_str = self.retrieve_context(query)
36         return self.generate_completion(query, context_str)
37
38
```

Results: This experiment yielded more poor results than Experiment 1. Further experimentation was necessary to achieve the desired performance.

Python ▾ as cell38

```
1 session.get_leaderboard()
```

app_name	app_version	Answer Relevance	Context Relevance	Groundedness	latency	total_cost
BreastCareTrial	simple	1	0.5	1	10.4623	1.6566
BreastCareTrial Retriever	base	None	1	None	0.273	0

Experiment 3:

Settings:

- Chunk size: 1180
- Chunk overlap: 350
- Retrieval limit: 3 (previously 4)

```
12
13 class text_chunker:
14
15     def process(self, pdf_text: str):
16         # Adjusted chunk size and overlap
17         text_splitter = RecursiveCharacterTextSplitter(
18             chunk_size = 1180, # Lower chunk size
19             chunk_overlap = 350, # Moderate overlap to maintain context
20             length_function = len
21         )
22
23         chunks = text_splitter.split_text(pdf_text)
24         df = pd.DataFrame(chunks, columns=['chunks'])
25
26         yield from df.iteruples(index=False, name=None)
27
28 $$$
```

Prompt:

```
1 from trulens_apps.custom import instrument
2 from snowflake.cortex import Complete
3
4
5 class RAG:
6
7     def __init__(self):
8         self.retriever = CortexSearchRetriever(snowpark_session=snowpark_session, limit_to_retrieve=3)
9
10    @instrument
11    def retrieve_context(self, query: str) -> list:
12        """
13        Retrieve relevant text from vector store.
14        """
15        return self.retriever.retrieve(query)
16
17    @instrument
18    def generate_completion(self, query: str, context_str: list) -> str:
19        """
20        Generate answer from context.
21        """
22        prompt = f"""
23        You are an expert assistant specializing in breast cancer clinical trials. Provide accurate answers strictly f
24        - Trial objectives, phases, and descriptions.
25        - Eligibility criteria and exclusion details.
26        - Trial locations, investigator contacts.
27
28        If the context does not contain the required information, respond with:
29        "I'm sorry, the provided context does not contain the information for your query."
30        Context: {context_str}
31        Question:
32        {query}
33        Answer:
34        """
35        return Complete("mistral-large2", prompt)
```

Results: By limiting the retrieved context to three chunks, this experiment achieved noticeable reductions in latency and cost. However, the relevance could still be improved.

Python v as cell38

```
1 session.get_leaderboard()
```

app_name	app_version	Answer Relevance	Context Relevance	Groundedness	latency	total_cost
BreastCareTrial	simple	1	0.8411	0.9992	3.9545	0.4944
BreastCareTrial Retriever	base	None	1	None	0.2376	0

Experiment 4:

Settings:

- Chunk size: 1200
- Chunk overlap: 350
- Retrieval limit: 3

```
14
15     def process(self, pdf_text: str):
16         # Adjusted chunk size and overlap
17         text_splitter = RecursiveCharacterTextSplitter(
18             chunk_size = 1200, # Lower chunk size
19             chunk_overlap = 350, # Moderate overlap to maintain context
20             length_function = len
21         )
22
```

Prompt:

```
6
7     def __init__(self):
8         self.retriever = CortexSearchRetriever(snowpark_session=snowpark_session, limit_to_retrieve=3)
9
10    @instrument
11    def retrieve_context(self, query: str) -> list:
12        """
13        Retrieve relevant text from vector store.
14        """
15        return self.retriever.retrieve(query)
16
17    @instrument
18    def generate_completion(self, query: str, context_str: list) -> str:
19        """
20        Generate answer from context.
21        """
22        prompt = f"""
23        You are an intelligent assistant specialized in breast cancer clinical trials.
24        |Your responses should focus on trial information, eligibility requirements, and next steps.
25        Context: {context_str}
26        Question:
27        {query}
28        Answer:
29        """
30        return Complete("mistral-large2", prompt)
31
```

Results: This experiment provided the best balance of:

- **Context Relevance**
- **Latency**
- **Cost Efficiency**


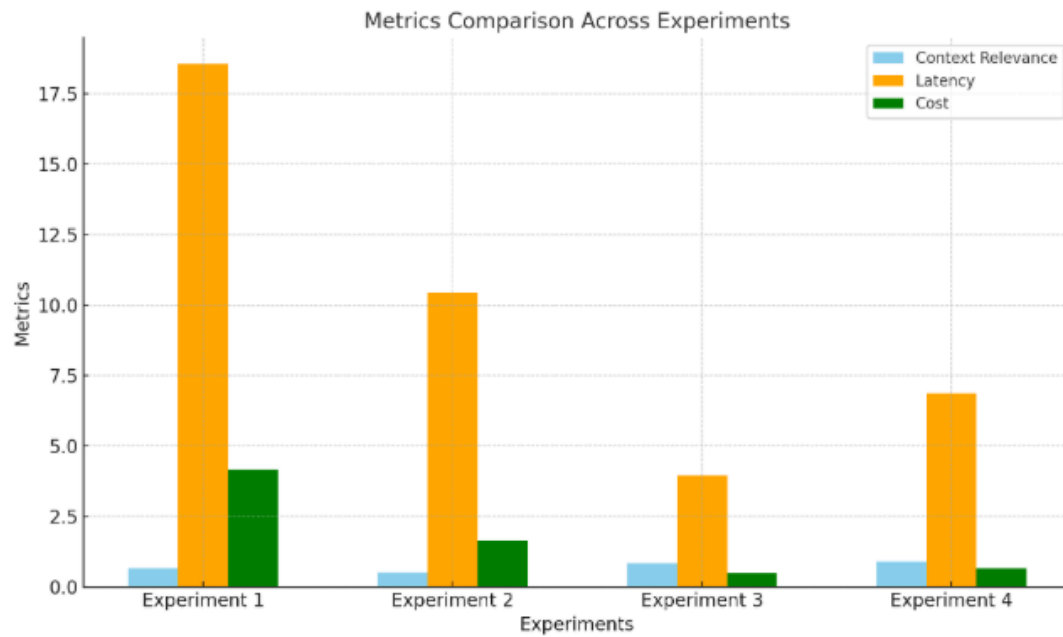
Experiment 4 is the most optimal configuration achieved during this analysis.

Python ▾ as cell38 0.0s ▶ ☰ ⋮

1 session.get_leaderboard()

app_name	app_version	Answer Relevance	Context Relevance	latency	total_cost
BreastCareTrial	simple	1	0.8889	6.8685	0.666
BreastCareTrial Retriever	base	None	1	0.2945	0





Conclusion

Based on the metrics analyzed, Experiment 3 and Experiment 4 show the most promising results.

- **Experiment 3** demonstrates a strong balance between high context relevance (0.84), low latency (3.95 seconds), and a minimal cost (0.49), making it highly efficient for performance optimization.
- **Experiment 4** further improves context relevance slightly (0.89) while maintaining reasonable latency (6.87 seconds) and cost (0.66).